


# Linear Regression

By Trần Minh Dương - Learning Support

## Overview

Imagine you're trying to predict house prices based on features like the size of the house, the number of bedrooms, and the location. Linear regression helps you establish a relationship between these features (inputs) and the price (output) by finding the best-fit line that minimizes the prediction error.



Linear regression offers a simple yet powerful way to model relationships in data. To find this best-fit line, we use one of two methods:

1. **Gradient Descent:** An iterative optimization process that adjusts the model step by step.
2. **Normal Equation:** A direct solution using linear algebra to calculate the optimal parameters.

In this document, we'll explore both methods in detail and understand their applications.

## 1. Gradient Descent

Gradient descent is an iterative optimization algorithm used to minimize the cost function in linear regression.

### Key Points:

**Hypothesis Function:** The predicted output of linear regression is computed as follows:

- **For a single example:**

$$h_{\theta}(x^{(i)}) = \theta^T x^{(i)} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$

Here,  $x^{(i)}$  is a single input example.

- **For the entire dataset (input matrix):**

$$h_{\theta}(X) = X\theta$$

Where:

- $x^{(i)}$ : A column vector representing the  $i$ -th example, containing  $n$  features.  
**Dimension:**  $n \times 1$ .
- $X$ : The input matrix containing  $m$  examples and  $n$  features.  
**Dimension:**  $m \times n$ .
- $\theta$ : The parameter vector containing  $n$  weights (including  $\theta_0$ , the bias term).  
**Dimension:**  $n \times 1$ .

**Goal:** Minimize the cost function (Mean Squared Error) to find the optimal parameters  $\theta$ .

**Mean Squared Error (MSE) Formula:**

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

Where:

- $m$ : Number of training examples in the batch.
- $h_{\theta}(x^{(i)})$ : Predicted output for the  $i$ -th example.
- $y^{(i)}$ : Actual output for the  $i$ -th example.

**Steps for Batch Gradient Descent of Linear Regression:**

1. Start with initial values for the parameters  $\theta$  (e.g.,  $\theta_0 = 0, \theta_1 = 0$ ).
2. Compute the gradient of the cost function  $J(\theta)$  with respect to each parameter:

- **Scalar form:**

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) \cdot x_j^{(i)}$$

where:

- $h_{\theta}(x^{(i)})$  is the predicted output
  - $h_{\theta}(x^{(i)}) - y^{(i)}$  is the error of the prediction compared to the real output
  - $x_j^{(i)}$  is the  $j$ -th feature value of the  $i$ -th training example
- **Matrix form:**

$$\nabla J(\theta) = \frac{1}{m} X^T (X\theta - y)$$

where:

- $X$  is the input matrix of dimension  $m \times n$ .
- $\theta$  is the column vector of dimension  $n \times 1$ .
- $X\theta$  gives the vector of predicted outputs of dimension  $m \times 1$ .

- $y$  is the vector of target values of dimension  $m \times 1$ .

3. Update each parameter using the formula:

$$\theta_j = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j} \quad \text{or} \quad \theta = \theta - \alpha \nabla J(\theta)$$

where  $\alpha$  is the learning rate.

4. Repeat until the cost function converges.

#### Advantages:

- Works well with large datasets.
- Scales to high-dimensional data.

#### Disadvantages:

- Requires choosing a suitable learning rate  $\alpha$ .
  - May converge slowly or get stuck in local minima if poorly initialized.
- 

## 2. Normal Equation

The normal equation is a closed-form solution for linear regression. It directly computes the optimal parameters  $\theta$  without iteration.

### Key Points:

**Goal:** Minimize the cost function by solving:

$$\nabla J(\theta) = 0$$

$$X^T(X\theta - y) = 0$$

$$X^T X \theta = X^T y$$

- Finally, we derive the formula for  $\theta$ :

$$\theta = (X^T X)^{-1} X^T y$$

where:

- $X$ : Feature matrix.
- $y$ : Vector of target values.
- $X^T$ : Transpose of  $X$ .
- $(X^T X)^{-1}$ : Inverse of  $X^T X$ .

#### Steps:

1. Construct the feature matrix  $X$  (including the column of 1s for the bias term).
2. Compute  $\theta$  using the formula above.

#### Advantages:

- No need for choosing a learning rate.
- Direct computation provides exact results.

#### Disadvantages:

- Computationally expensive for large datasets due to matrix inversion ( $O(n^3)$ ).
- May cause error when  $X^T X$  is non-invertible or the dataset is too large to fit into memory.

## Summary of Key Differences

Aspect	Gradient Descent	Normal Equation
Approach	Iterative optimization	Direct computation
Computation Cost	Scales well for large datasets	Expensive for large datasets
Parameters Required	Requires learning rate ( $\alpha$ )	No hyperparameters
Suitability	Large datasets or high dimensions	Small to medium datasets

## Exercise

Given the training data for a linear regression problem, to predict the final score of a student in Machine Learning class, based on his/her score in programming and probability class. Starting at:

- $\theta_0 = 0$ ,  $\theta_1 = 1$ ,  $\theta_2 = -1$
- Learning rate ( $\alpha$ ) = 4

### Task:

1. Calculate the coefficients after the first iteration with **batch-gradient descent**.
2. Based on the coefficients you just found, calculate the prediction score of the fifth student.

### Data Table:

Student	Score in Programming	Score in Probability	Score in Machine Learning
1	18	15	15
2	15	10	11
3	16	12	13
4	19	16	17
5	17	11	???

```
In [2]: #In this exercise, I will use the matrix approach with the numpy library in python
import numpy as np
```

```
In [3]: #Input / Output
X = np.array([[18,15],[15,10],[16,12],[19,16]])
y = np.array([15,11,13,17])

#Parameters
```

```
theta = np.array([0,1,-1])
alpha = 4
m = len(y) #batch size
```

```
In [4]: #reshape X to add an additional column of 1 to the left (bias term)
X_bias = np.hstack((np.ones((m,1)),X))
print("X_bias = \n",X_bias)
```

```
X_bias =
[[ 1. 18. 15.]
 [ 1. 15. 10.]
 [ 1. 16. 12.]
 [ 1. 19. 16.]]
```

```
In [5]: def hypothesis(X,theta):
        return X @ theta # the @ here is the operator for matrix multiplication in python

predictions = hypothesis(X_bias,theta)
print(predictions)
```

```
[3. 5. 4. 3.]
```

```
In [6]: errors = predictions - y
print(errors)
```

```
[-12. -6. -9. -14.]
```

```
In [7]: gradient = 1/m * (X_bias.T @ errors)
print(gradient)
```

```
[ -10.25 -179.   -143.   ]
```

```
In [8]: theta = theta - alpha * gradient
print("Theta after one iteration of batch-gradient descent is: \n",theta)
```

```
Theta after one iteration of batch-gradient descent is:
[ 41. 717. 571.]
```

```
In [9]: fifthStudent = np.array([1,17,11])

score = hypothesis(fifthStudent,theta)
print("The score of the fifth student with 17 in Programming and 11 in Probability is
```

```
The score of the fifth student with 17 in Programming and 11 in Probability is: 18511.0
```

---

## Normal equation approach:

```
In [10]: X_bias = np.hstack((np.ones((m,1)),X))

# Compute theta = (X^T . X)^(-1) . X^T . y
# np.linalg.inv(A) returns the inverse of A : A^-1
theta_normal = np.linalg.inv(X_bias.T @ X_bias) @ X_bias.T @ y
print("Theta from the normal equation:", np.round(theta_normal,2))
```

```
Theta from the normal equation: [-9.8  1.4 -0. ]
```

---

This document was created in Jupyter Notebook by [Trần Minh Dương \(tmd\)](#).

If you have any questions or notice any errors, feel free to reach out via Discord at [@tmdhoctiengphap](#) or [@ICT-Supporters](#) on the USTH Learning Support server.

Check out my GitHub repository for more projects: [GalaxyAnnihilator/MachineLearning](#) .